



MSKCC

Automatic Nuclei Segmentation

for Histopathology Images

AI Studio Final Presentation

Break Through Tech New York @Cornell Tech
December 16 2022



Introductions



Our Team



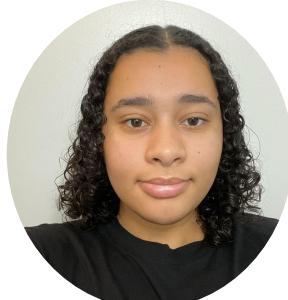
Chichi Mba
Columbia University



Nicole Joseph
The Cooper Union



Kaylee Fang
Columbia University



Leona Torres
Queens College



Selina Sun
New York University



Our AI Studio TA and Challenge Advisor



Ananya Ganesh
AI Studio TA



Alex Hollingsworth
Challenge Advisor



Presentation Agenda

1. Introductions
2. Project Overview
3. Data Preprocessing
4. Model Selection and Evaluation
5. Final Thoughts



AI Studio Project Overview



“

Create a deep learning model that can
accurately segment nuclei and separate
overlapping or touching nuclei in H&E
stained tissue sections



Our Goals

- Learn more about deep learning, computer vision, and image processing techniques
- Understand how to work with a small data set
- Differentiate between the various types of image segmentation
- Create a neural network that can accurately segment nuclei and separate overlapping and touching cells in H&E-stained tissue sections using Deep Learning.

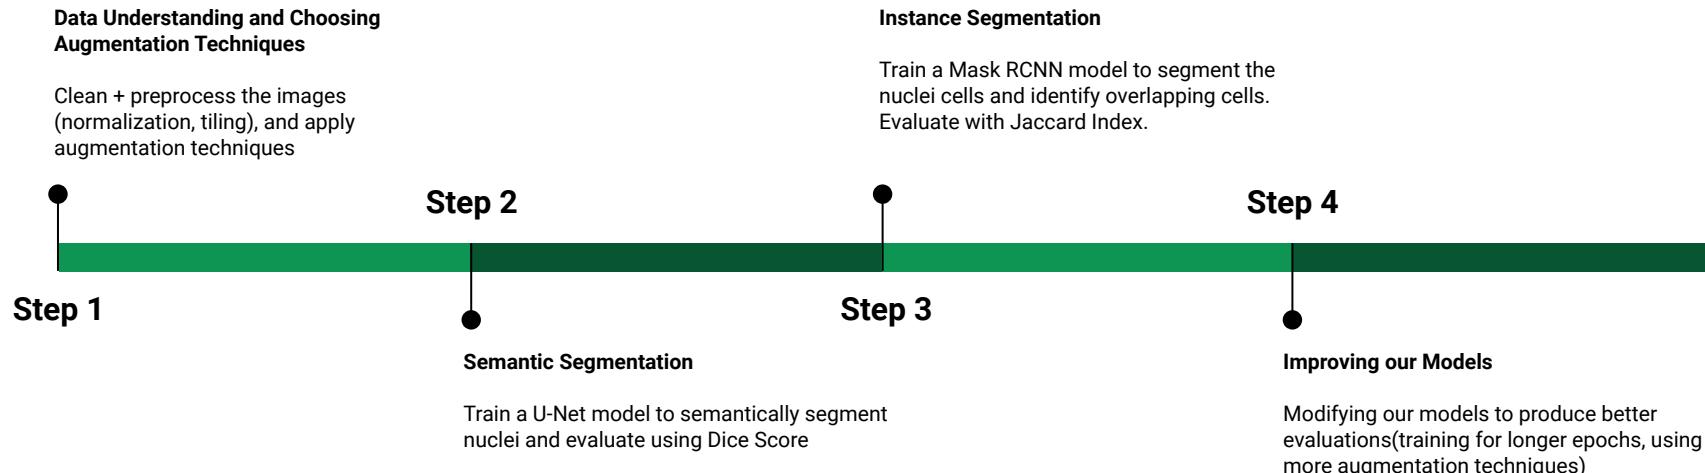


Business Impact

- Accurately segmented nuclei provide important insights for pathologists:
 - Physical characteristics of nuclei (size, shape)
 - Spatial distribution
- Reduces need for manual annotations of nuclei
- Manual separation of images is tedious and time consuming
- Automated nuclei segmentation is important to gaining further insight into cell features and functionality



Our Approach





Resources We Leveraged



TensorFlow



matplotlib



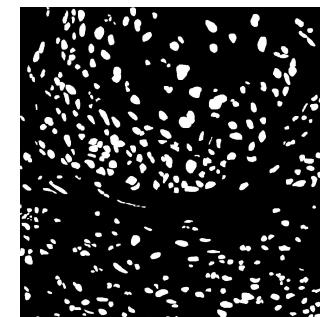
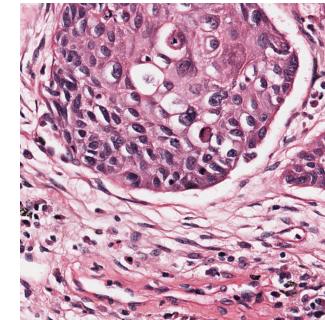
Data Preprocessing



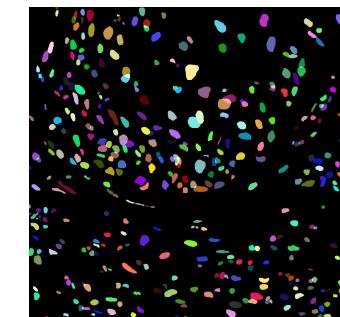
The Dataset

- **MoNuSeg dataset**
 - Public domain data
 - Obtained by annotating tissue images of several patients with tumors of different organs at multiple hospitals
 - Tissue images across 30 patients and 7 different organs
- 44 H&E stained tissue images, each 1000x1000
 - 30 for training, 14 for testing

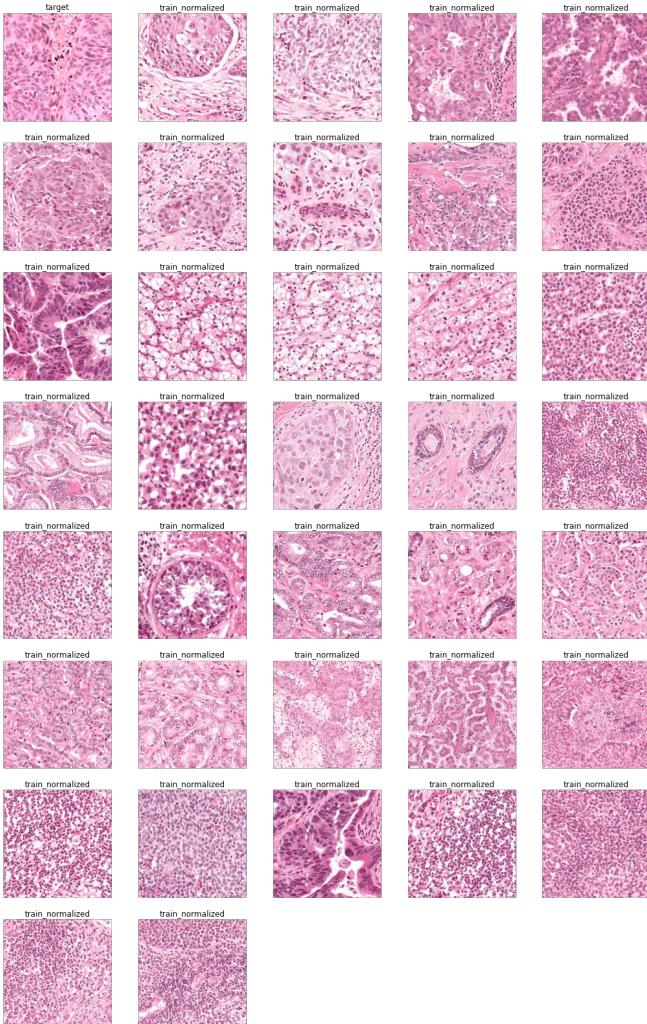
H&E stained tissue



Ground-Truth Binary Mask



Ground-Truth Color Mask



Stain Normalization

- Every image in the training and test dataset was normalized according to one random image from the dataset
 - Normalizes based on stain color and hue
 - Transforms all features to be on a similar scale
- Improves model training and performance
- Tools used: **StainTools** library
(<https://staintools.readthedocs.io/en/latest/>)



Image Tiling

- Each 1000x1000 image + mask was tiled into 16 256x256 tiles
- Benefits of tiling:
 - Faster training and performance
 - Increases batch diversity
- Tools used: **tiler** library
(<https://pypi.org/project/tiler/>)

▼ Tiling

```
✓ [55] # tiling tissue images
       tile_dataset(data)
```





Preprocessing Object

```
[ ] class PreProcessing:
    def __init__(self, path_to_training, path_to_testing):
        self.training_imgs_path = path_to_training
        self.testing_imgs_path = path_to_testing

        # names of images
        self.train_names = np.sort(np.array([file_name for file_name in.listdir(self.training_imgs_path) if isfile(join(self.training_imgs_path, file_name))]))
        self.test_names = np.sort(np.array([file_name for file_name in.listdir(self.testing_imgs_path) if isfile(join(self.testing_imgs_path, file_name))]))

        # normalize images
        self.training_imgs = []
        self.testing_imgs = []

        # normalized and tiled images
        self.training_imgs_tiled = []
        self.testing_imgs_tiled = []

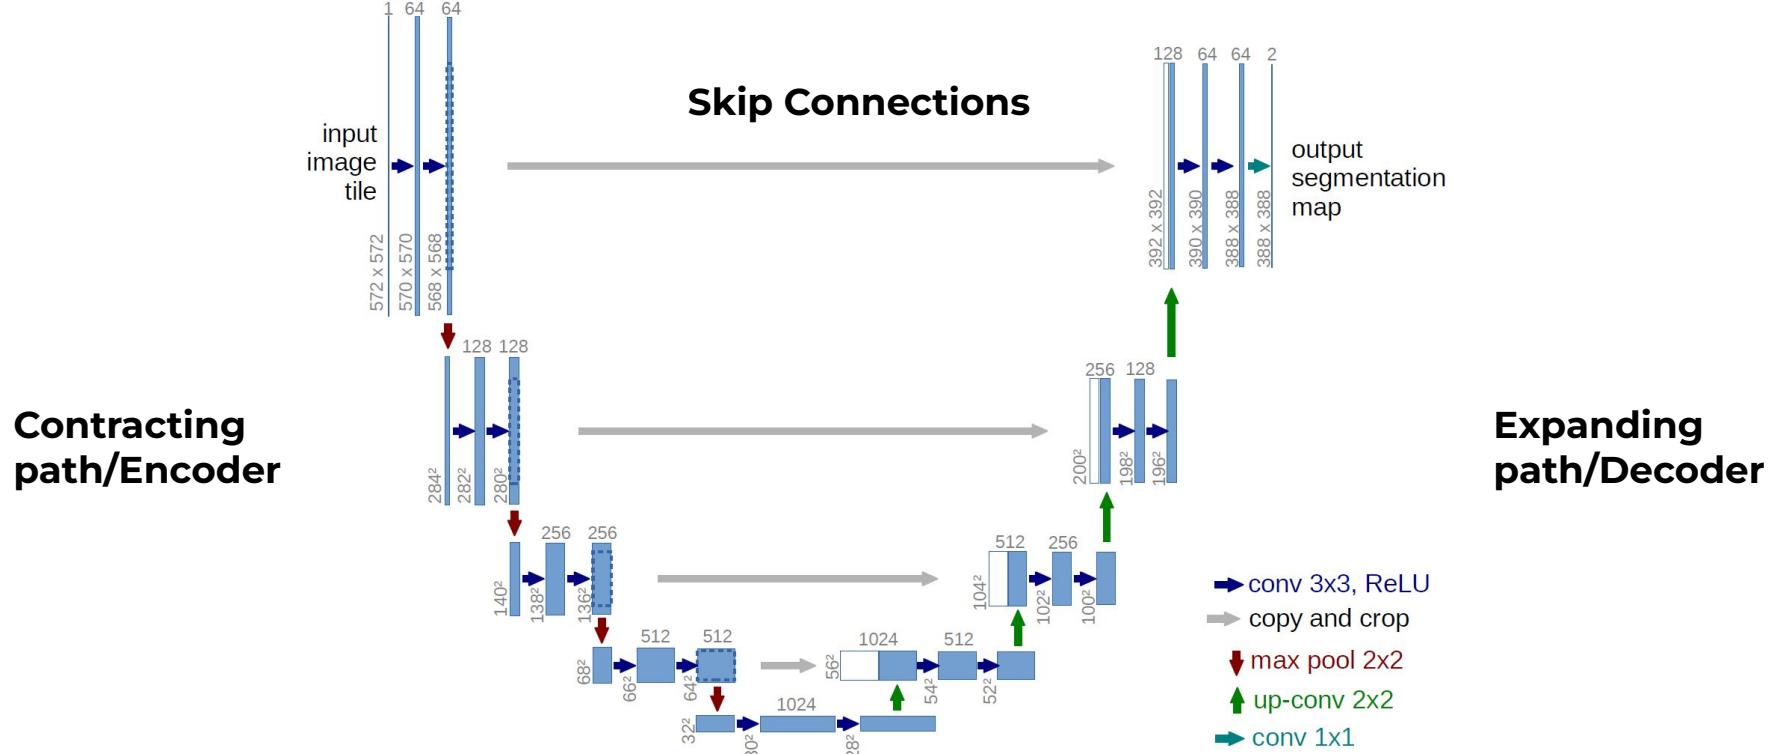
        # normalized, tiled and augmented training images
        self.training_imgs_aug = []
```



Model Selection and Evaluation



U-Net Architecture



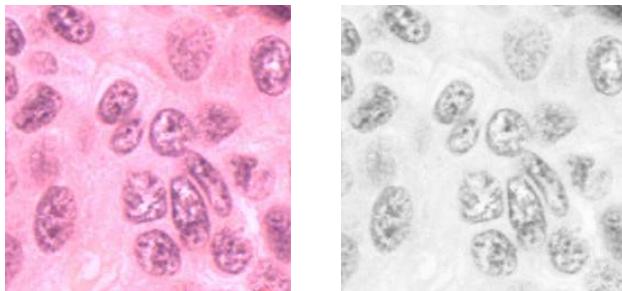
Example Code

```
#Contraction/Encoder path (going down the U) - this is made up of convolutional layers (c1, c2, etc)
# feature space = 16; apply 3x3 convolutions; padding = same to keep output width and height same as input
c1 = Conv2D(16, (3, 3), activation='relu', kernel_initializer='he_normal', padding='same')(s)
c1 = Dropout(0.1)(c1)
c1 = Conv2D(16, (3, 3), activation='relu', kernel_initializer='he_normal', padding='same')(c1)
p1 = MaxPooling2D((2, 2))(c1)
```

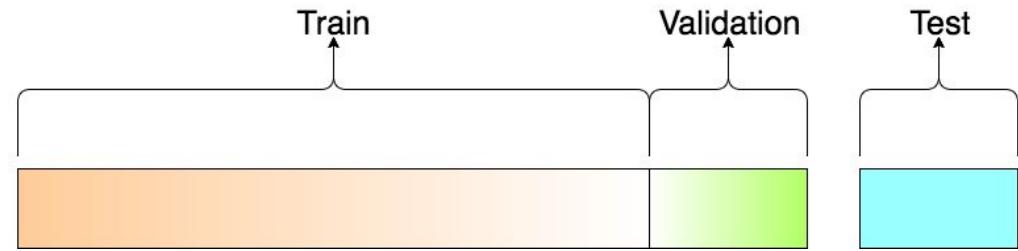
- Scaled down feature space dimensions by a factor of 8
- For the outputs, we used a sigmoid activation function to ensure that the mask pixels are in a [0, 1] range



Training Process



Stage 1: When we first started training the U-Net model, we converted the images to grayscale
Best Accuracy Score: ~20%

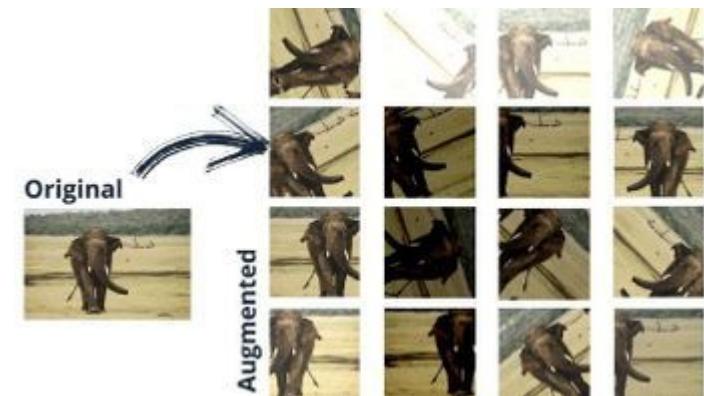


Stage 2: Added a validation set that was 10% of the training data
Best Accuracy Score: ~40%



Stage 3: Data Augmentation (on-the-fly)

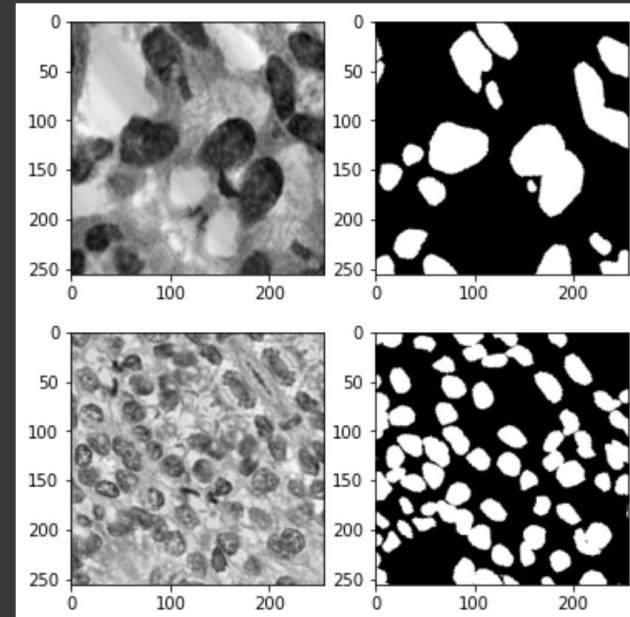
```
img_data_gen_args = dict(rotation_range=90,  
                         width_shift_range=0.3,  
                         height_shift_range=0.3,  
                         zoom_range=0.3,  
                         horizontal_flip=True,  
                         vertical_flip=True,  
                         fill_mode='reflect')
```



Stage 3: Data Augmentation increases the amount of data needed to train robust AI models. Applied to both images and masks.
Best Accuracy Score: ~60%

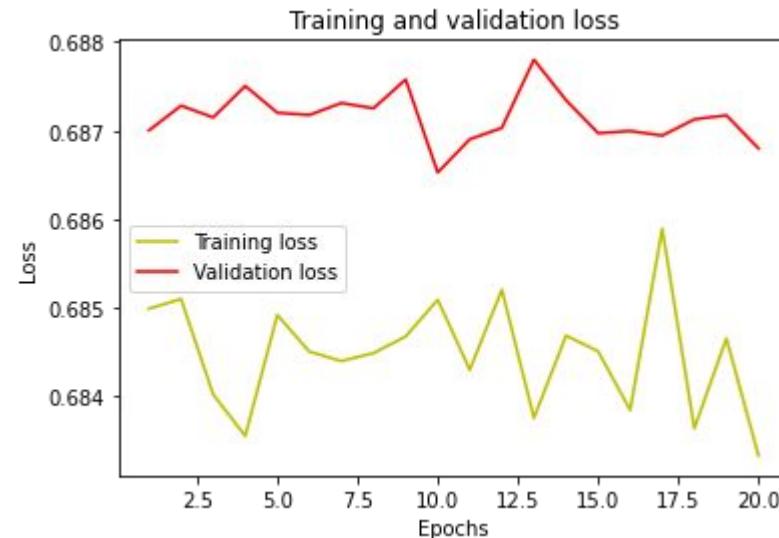
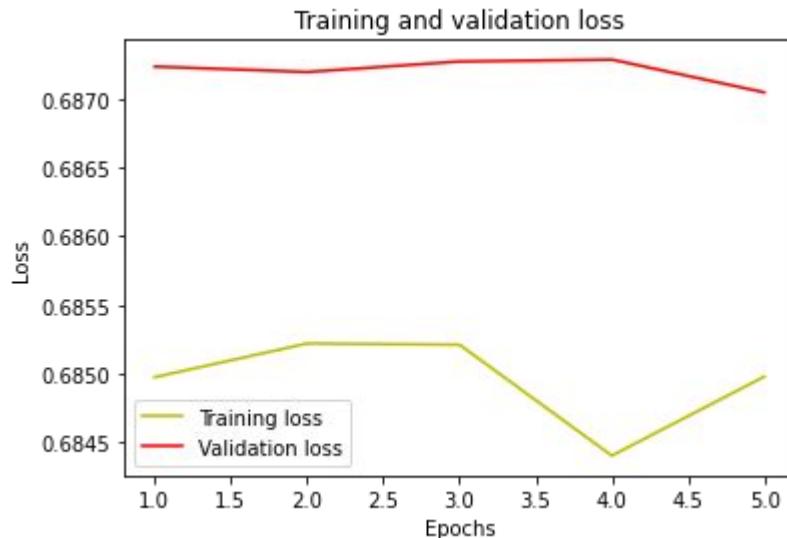
On-the-fly data augmentation via the keras Image Data Generator ensures that the augmented images are matched up with their corresponding masks

```
[57] x = image_generator.next()
y = mask_generator.next()
for i in range(0,2):
    image = x[i]
    mask = y[i]
    plt.subplot(1,2,1)
    plt.imshow(image[:, :, 0], cmap = 'gray')
    plt.subplot(1,2,2)
    plt.imshow(mask[:, :, 0], cmap = 'gray')
plt.show()
```



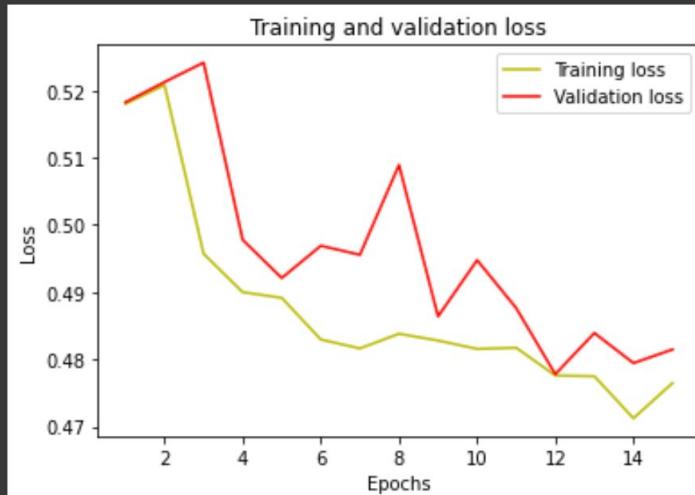


U-Net Results



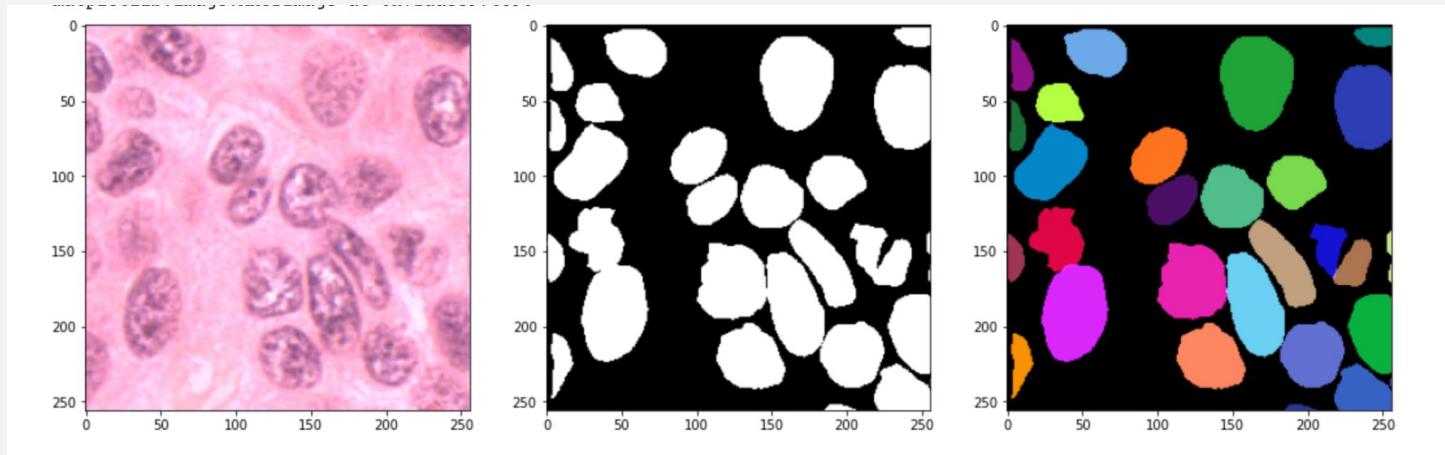
Stage 4: Fix mask and image tile error (some were not matched up)
Best Accuracy Score: ~75%

```
[64] loss = history.history['loss']
      val_loss = history.history['val_loss']
      epochs = range(1, len(loss) + 1)
      plt.plot(epochs, loss, 'y', label='Training loss')
      plt.plot(epochs, val_loss, 'r', label='Validation loss')
      plt.title('Training and validation loss')
      plt.xlabel('Epochs')
      plt.ylabel('Loss')
      plt.legend()
      plt.show()
```





Mask R-CNN for Instance Segmentation





Model Descriptions

Model Name	Description	Results	Why we chose it?
U-Net	Semantic segmentation - segments objects in image from background	A binary mask (black and white)	Gold standard for biomedical image segmentation problems
Mask R-CNN	Instance segmentation - combines image segmentation with object detection to segment objects from each other	A mask that shows each distinct nuclei cells	Proven to be a well performing model for instance segmentation



Final Thoughts



Obstacles

- We had to ensure that the dimensions of the input images for the U-Net matched up
- The mrcnn package for the Mask RCNN model is only compatible with a certain version of tensorflow which forced us to use virtual machines
- Using virtual machines on Google Cloud allowed us to run our U-Net model for more epochs to get a more accurate model



Insights and Key Findings

- On the fly augmentation removes the problem of having to match up masks and the predicted masks
- Converting the training images to grayscale may simplify computations and reduce time
 - However, training with RGB images provides information in all three channels and improves model performance
- The Dice Score and Jaccard Index are similar metrics for evaluating model performance, but the Dice Score provides more nuance in comparing pixel similarity



Potential Next Steps

- Get more training data
- Experiment with U-Net architecture
- Implement advanced techniques for hyperparameter tuning
 - Learning Rate Scheduler
- Watershed segmentation to improve accuracy



Questions?

